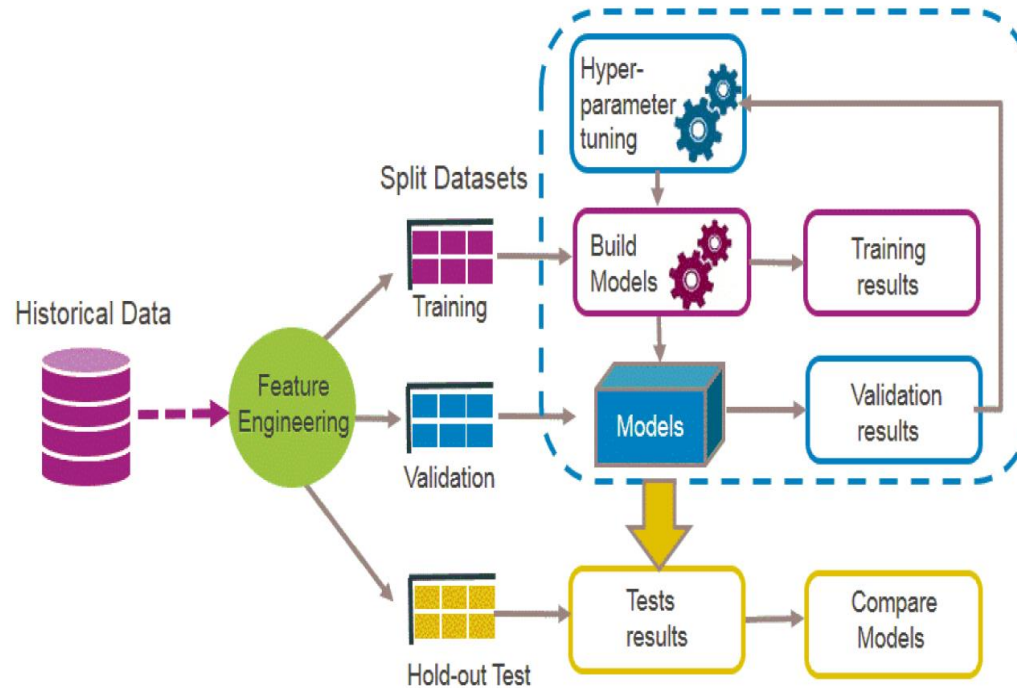
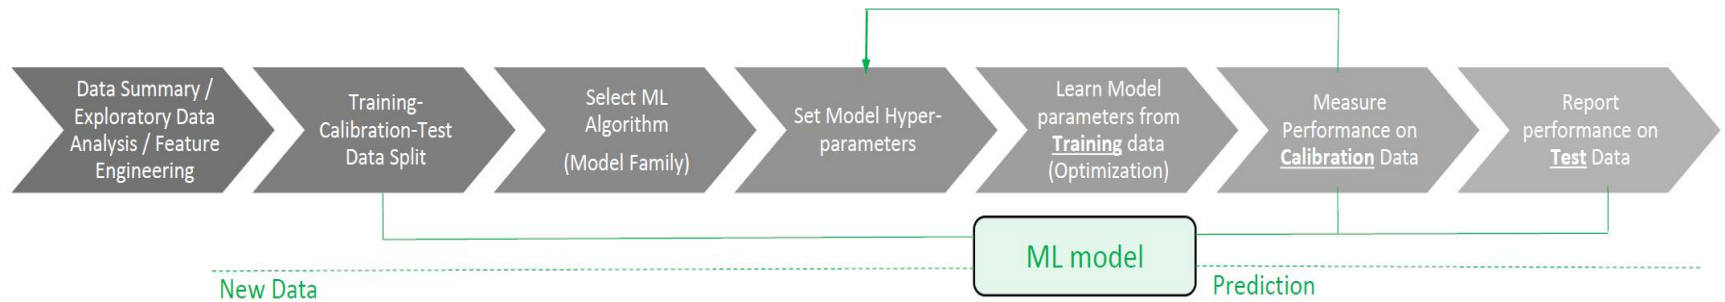


Predictive Modelling to determine Vehicle Prices

Regression



Problem Statement

Given the historical data, the aim is to build a predictive model to determine the price of used car vehicles in India.


- From the problem statement it is evident that it is a regression problem.
- Segments benefitted : Car manufacturers, car dealers, banks, individuals

Data Preprocessing

- Changed the data types to numeric, factors, and Strings as per given data description.
- Removed the Data columns
- Removed the NameOfTheVehicle String column
- Removed the columns exhibiting NearZero Variance
- Removed the ID and Zipcode columns

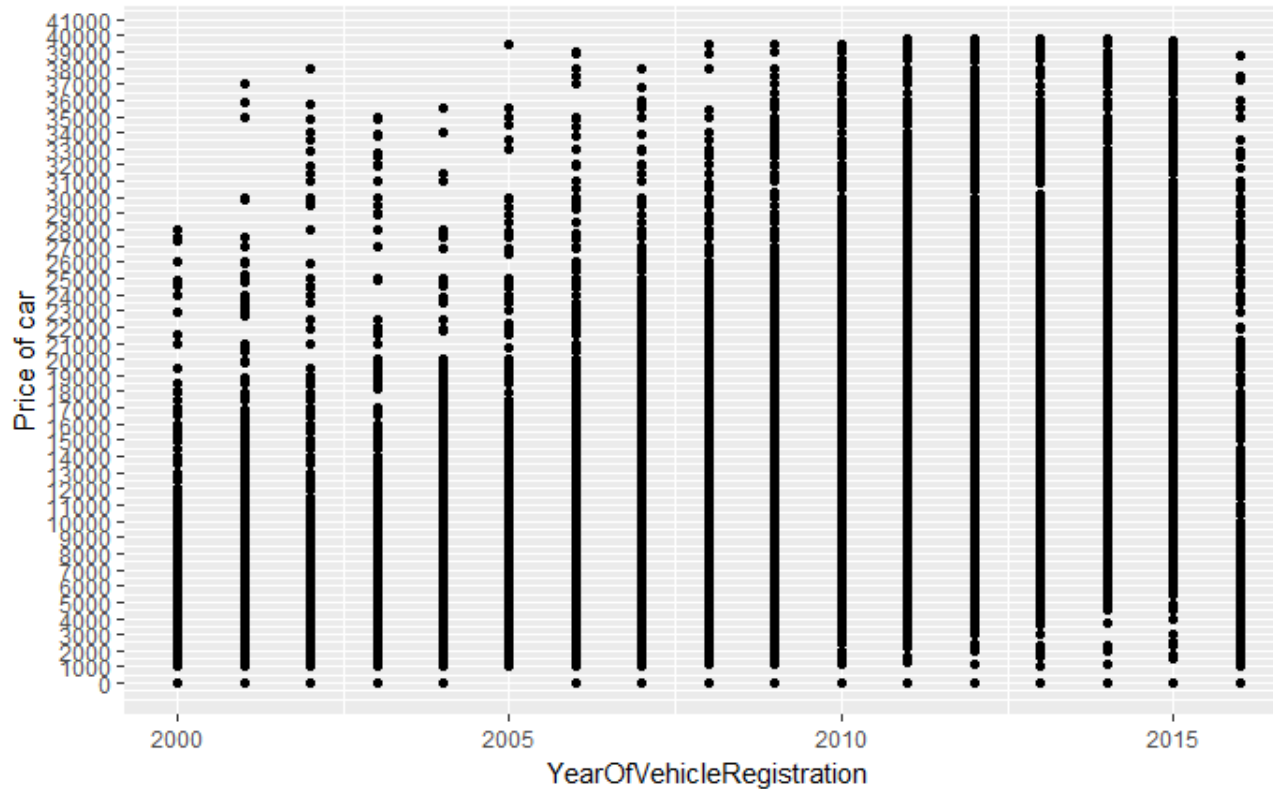
Data Exploration

- Most of the Engine Powers associated with higher power don't have higher price range. → Outliers




Data Exploration

- Older the Vehicle Registration, lower the price



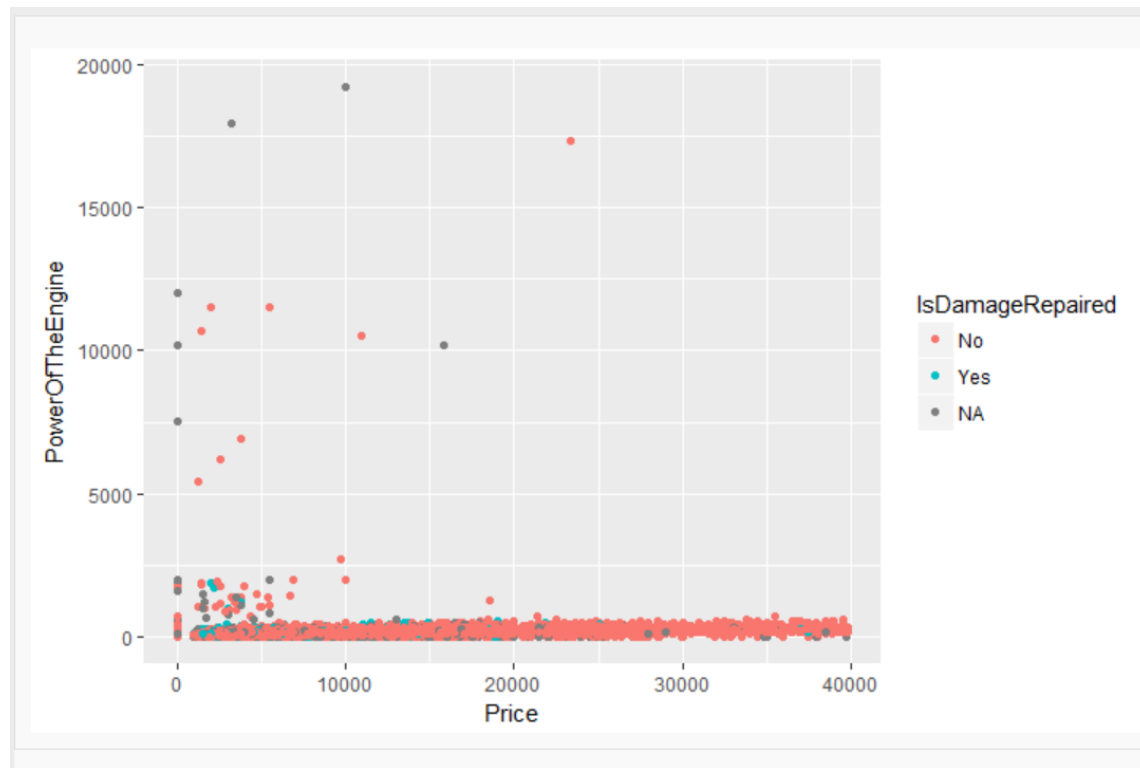
Data Exploration

- ModelOfTheVehicle Frequency distribution. => Decided to make factors, and combine those with frequency < 3% to reduce the levels.



Data Exploration

- Very few observations with IsDamageRepaired as “Yes” => But decided to keep this attribute



Data Cleaning

- Dataset (train + test) split into numeric and categorical separately.
- Checked for numeric predictors having high correlations (>0.80)
- No missing value was found in numerical attributes
- Created a new factor for missing categorical attributes
- Combine categorical variables which have several levels with low frequencies ($\leq 3\%$) and created new factor “Other”
- Standardized the numeric data

Preparation of Train/Test Data



- Test data : New Test data
- Train/Valid Data : Combined the old test data and original training data, and randomly split the data into ratio 70-30 using library caTools
- Created entire data into numeric format, using `model.matrix`
- Final Dimesions : 44

Model Intuition

- I decided to adopt a workflow testing several different models and check the respective error metrics.
 - Models Built :
 - Linear Regression
 - CART
 - Random Forest
 - XGBoost
- After that I used stacking to check if model MAPE improves.


Linear Regression

- The model was firstly built and then updated by removing observations based on cooks distance.
- The linearity assumptions didn't hold valid.



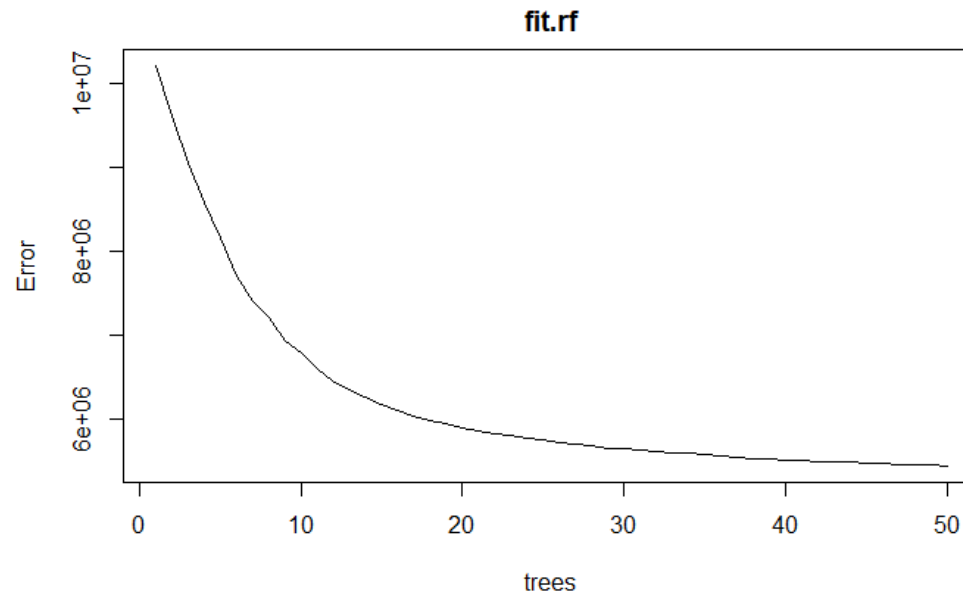
CART

- Built and tuned the `cp` parameter to prune the tree



Random Forest

- Tuned the ntree and mtry parameters to obtain optimal metrics.
- % Var explained: 88.09



Extreme Gradient Boosting using DT

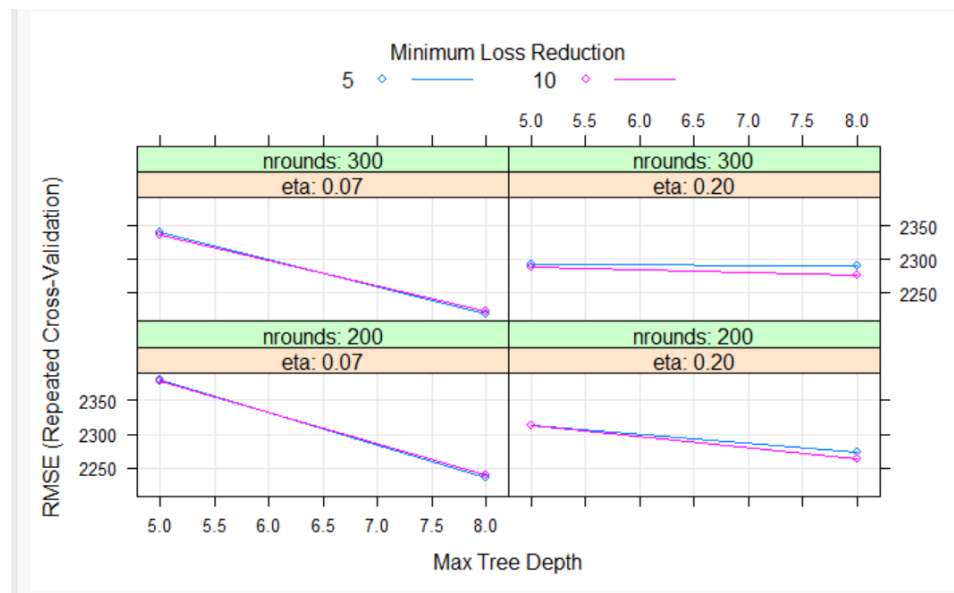
- Firstly did a random grid search, and then manually gave range of hyperparameters for model to find best combination.
- The later gave better results.

Tuning parameters:

- `nrounds` (# Boosting Iterations)
- `max_depth` (Max Tree Depth)
- `eta` (Shrinkage)
- `gamma` (Minimum Loss Reduction)
- `colsample_bytree` (Subsample Ratio of Columns)
- `min_child_weight` (Minimum Sum of Instance Weight)
- `subsample` (Subsample Percentage)

Extreme Gradient Boosting using DT

- Hyper parameter Tuning



- Best Tune

	nrounds <dbl>	max_depth <dbl>	eta <dbl>	gamma <dbl>	colsample_bytree <dbl>	min_child_weight <dbl>	subsample <dbl>
6	300	8	0.07	5	0.8	1	0.8

Stacking

- After that I employed Stacking to combine the results.
- The correlations in stacked predictors were quite high. So I did not expect a significant improvement after stacking.
- However I gave it a try, and see if error metrics improve. Models tried as meta learners : RF, Linear Regression and CART

Stacking

- Correlations


	RF	XGBTREE	CART	Price
RF	1.0000000	0.9897701	0.8983955	0.9854124
XGBTREE	0.9897701	1.0000000	0.9030494	0.9701990
CART	0.8983955	0.9030494	1.0000000	0.8610358
Price	0.9854124	0.9701990	0.8610358	1.0000000

- Head of data

	RF <dbl>	XGBTREE <dbl>	CART <dbl>	Price <dbl>
1	3774.896	4063.092	2719.588	3850
3	6568.484	8208.192	6195.689	5990
4	4095.842	4552.498	6195.689	4000
5	11980.811	9968.633	12068.061	12950
6	3213.363	2623.744	2719.588	3600
7	3095.410	3322.637	2719.588	4450

6 rows


Stacking using Linear Regression



Adjusted R-squared: 0.9747

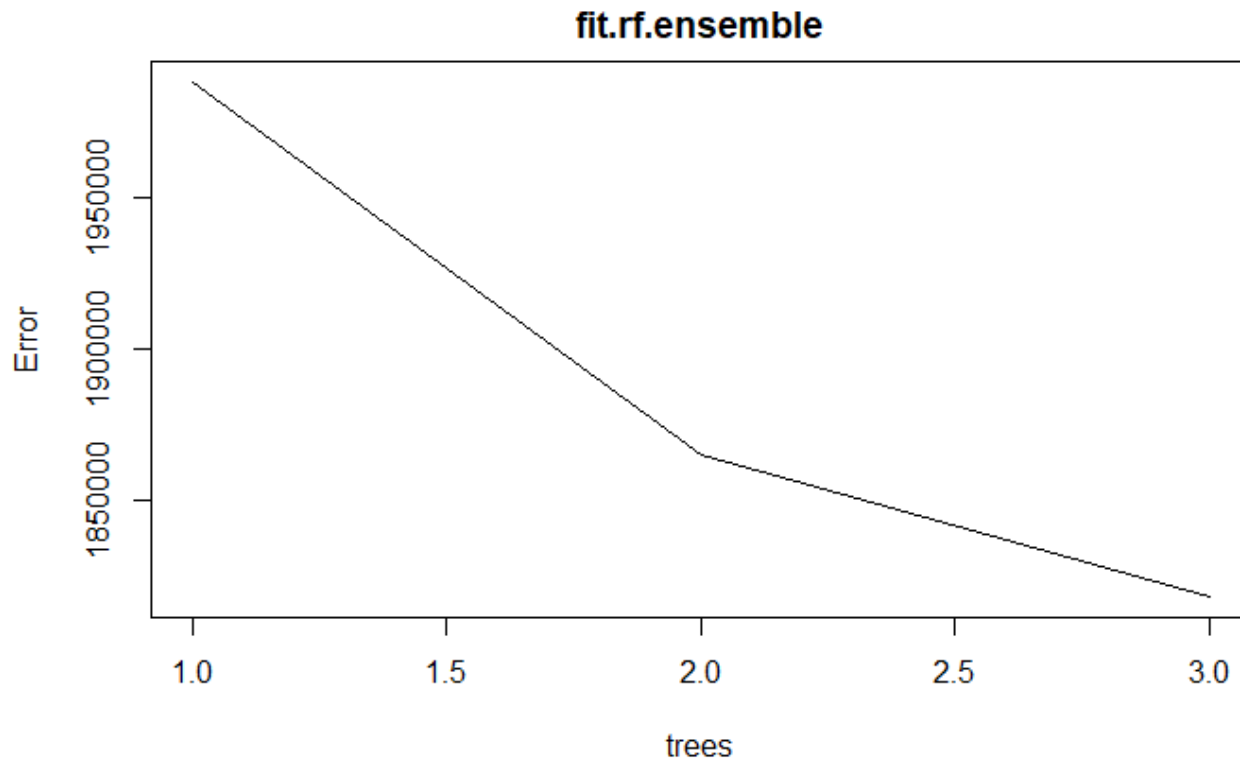
Stacking using CART

- Cp tuning



Stacking using Random Forest

- % Var explained: 96.02



Comparison of standalone models

MAPE	Train Error	Validation Error
Linear Regression	1.521916e+00	1.962097e+00
CART	4.424521e-01	4.469816e-01
Random Forests	1.264031e-01	2.476127e-01
XGBoost (Random Grid Search)	2.054038e-01	2.397260e-01
XGBoost (Manual Grid Search)	2.612451e-01	2.657679e-01
Stacking using CART	1.991522e-01	2.854830e-01
Stacking using Linear Regression	1.140418e-01	2.535492e-01
Stacking using RF	1.060275e-01	2.532537e-01

Results

- I got best results using Random Forest standalone model.
- Stacking improved the train error, but test error didn't decrease
- Least Test Data MAPE came approximately 39%

Your answer passed the tests! Your score is 2.58%

Congratulations!! Your model surpassed the baseline model with MAPE of 38.8233%

However your score will be scaled post submission deadline.

Further Improvements

- I would like to later work on this data to improve the model. Some of the ideas I have are
 - Clean the “NameOfTheVehicle” attribute using string regex techniques and add into model
 - Use the date columns in the model